

Group-BeyondAI-Q3-Proposal-1

November 15, 2020

1 EE4211 Data Science for IoT Project

1.1 Group Name: BeyondAI

Group Members: Chen Yiming (e0622334@u.nus.edu) Matric No: A0226328E
LIU XINGYU (e0506559@u.nus.edu) Matric No: A0116430W
YEAP CHUN LIK (e0319167@u.nus.edu) Matric No: A0185876H
ZHANG GUYU (e0506558@u.nus.edu) Matric No: A0116427J

2 Instructions:

In this project, you are given a dataset collected by an actual IoT system (see description below) and asked to use the dataset to build a forecasting model. You have to answer a set of questions, as well as propose your own interesting questions. 1. Form teams in groups of 4 students and select a name for your team. Be creative! Please email me your group members and team name. 2. Complete Question 1. Use a Jupyter notebook (ipynb file) to do the analysis and answer all the parts of the Question. Submit (i) PDF file/Print preview of your Jupyter notebook, and (ii) the Jupyter notebook (ipynb file). Zip both files into one zip file named as GroupName Question 1.zip and upload it to the appropriate LumiNUS folder. 3. Do the same for Question 2. Please name your file GroupName Question 2.zip and upload to LumiNUS. 4. Do the same for Question 3. For Question 3, please include a detailed description of your proposed work. Please name your file GroupName Question 3.zip and upload to LumiNUS. 5. The project carries a total of 40 marks: 30 marks for technical contributions (10 marks for each question), and 10 marks for presentation.

3 Data File:

The data file is available in the IVLE workbin under the directory "Project Details".

4 Data Description:

In this project, we will consider natural gas consumption data from residential consumers. The smart gas meter data used for this paper was obtained from the Pecan Street project (<https://www.pecanstreet.org/>). The source of the data are homes in the Mueller neighborhood of Austin, Texas, USA. The homes in this neighborhood are primarily newly constructed, and include single-family homes, apartments, and town homes. Itron Centron SR smart gas meters are deployed in these homes and these meters send their information to a gateway inside the home. The gateway uses the home's Internet connection to send the data to the meter data management

system (MDMS) or the processing center. The gas meters measure the cumulative gas consumption at a frequency of 15 seconds. The meters report a reading (in terms of the cumulative consumption) when the last marginal 2 cubic foot (or higher) of natural gas passes through the meter. Data from a six month interval (1 Oct 2015 to 31 Mar 2016) has been provided. The data has the following format: <Timestamp (localtime)> <MeterID (dataid)> <meter reading (meter_value)> The timestamp provides the date as well as the the hour and minute values when each reading was taken. Each meter has an unique identifier (MeterID). Recall that the meter readings are cumulative and not generated at periodic intervals.

5 Additional Information about Data Collection:

1. Gas flow meters have a sensor that is used to measure the volume of gas that passes through a pipe. Different meters use different sensors (e.g. ultrasonic sensors, synthetic diaphragm with rotating valve etc.). The meters check on the sensors periodically to get a reading of the current consumption value. This is what is meant in the sentence above: "The gas meters measure the cumulative gas consumption at a frequency of 15 seconds."
2. Now, just because the meter has obtained a reading from the sensors, it does not have to send the reading off to the meter data management system (MDMS). Imagine 1.3 million households in Singapore sending out gas readings every 15 seconds to Singapore Power. The processing and bandwidth requirements may be too high for Singapore Power. So Singapore Power may wish for the meters to report at a lower frequency or when the consumption exceeds a certain threshold. However, the smart meter manufacturer does not know what is the reporting criterion of its users. So it builds meters that can read every 15 seconds because it thinks that this is a frequency that is high enough for all potential customers. The "reporting" frequency to the MDMS (as opposed to the "measuring" frequency) can be determined by the user of the meter such as Singapore Power.
3. So when are the meters supposed to "report" to the MDMS? The documentation that came with the data says "once the marginal consumption exceed 2 cubic meters". As you may observe in the data, this is not necessarily the case in some of the readings. So is that an anomaly? That is for you to decide and justify. If you were Singapore Power, under what circumstances would you think that a meter reading is suspicious and decide to investigate? Remember that there are two sides to the story. If you do not receive a reading from a meter for a really long time, would you think that the meter is defective? So would that justify sending a reading even if the consumption has not increased?

6 Questions:

6.1 3. Student Proposal (10 marks)

6.2 Proposal 1

First load the data

```
[265]: # load necessary library

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

from datetime import datetime, date
from scipy import stats
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures

import warnings
warnings.filterwarnings('ignore')

```

[266]: *# read original data and take out the seconds and UTC offset*

```

data = pd.read_csv("dataport-export_gas_oct2015-mar2016.csv")
data['localminute'] = data['localminute'].astype(str).str[:19]
data['localminute'] = pd.to_datetime(data['localminute'])

```

[267]: *# check the first 5 rows in the dataset*

```
data.tail()
```

```

[267]:
      localminute  dataid  meter_value
1584818 2016-03-31 23:59:14      2129      201726
1584819 2016-03-31 23:59:17      2945      161232
1584820 2016-03-31 23:59:35      9729      138146
1584821 2016-03-31 23:59:47      5129      166488
1584822 2016-03-31 23:59:58       484      114174

```

[268]: *# check the data information of the dataset*

```

print(data.info())
print('-----')
print(data.describe(include = "all"))

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1584823 entries, 0 to 1584822
Data columns (total 3 columns):
localminute    1584823 non-null datetime64[ns]
dataid         1584823 non-null int64
meter_value    1584823 non-null int64
dtypes: datetime64[ns](1), int64(2)
memory usage: 36.3 MB
None
-----

```

	localminute	dataid	meter_value
count	1584823	1.584823e+06	1.584823e+06
unique	1499587	NaN	NaN
top	2015-12-05 20:57:14	NaN	NaN
freq	6	NaN	NaN
first	2015-10-01 00:00:10	NaN	NaN

last	2016-03-31 23:59:58	NaN	NaN
mean	NaN	4.352815e+03	2.015056e+05
std	NaN	2.941902e+03	1.351182e+05
min	NaN	3.500000e+01	2.829800e+04
25%	NaN	1.714000e+03	1.145800e+05
50%	NaN	4.031000e+03	1.670940e+05
75%	NaN	7.017000e+03	2.364940e+05
max	NaN	9.982000e+03	8.158240e+05

```
[269]: # check is there any null vales in the dataset

print("Number of null value is {}".format(data.isnull().sum()))
```

```
Number of null value is localminute    0
dataid          0
meter_value     0
dtype: int64
```

```
[270]: # create the list that contains all the meter ID

lstMeter = data['dataid'].unique()
```


```
[271]: # seperate all readings according to the meter ID

dictMeter = {c: pd.DataFrame(data[data['dataid']==c]) for c in lstMeter}
```

6.3 Discussion: preprocessing the raw data first


Same as Q2, firstly clean the data and remove the outliers

Below are an example for meter ID 3544:



We can see that there are some outliers malfunction readings during 2015-12 to 2016-01 and we should remove those outliers as it may affect our modeling.

Cleaned example for meter ID 3544:



[272]: *# clean the original data and remove thoes ooutliers*

```
dictMeter_clean = {}
for key in dictMeter:
    dictMeter_clean[key] = dictMeter[key].
    ↳drop(dictMeter[key][dictMeter[key]['meter_value']>dictMeter[key].\
        iloc[-1,:]['meter_value']].index)
```

[273]: *# divide the 6 months time into 24h * 183d = 4392 segments. For each meter, get*
↳the average meter reading of all the
readings within that segment. For example, segment 0 means period 2015-10-01
↳00:00:00 to 2015-10-01 01:00:00. So
the first reading is the average reading of all the readings in this period.
↳After processing, all the meter will
have the same length which is 4392 readings. If there is no reading for
↳certain meter for certain time period, the
reading will be nan.

```
hourly_readings = {}
for key in dictMeter_clean:
    start_time = pd.Timestamp(year=2015, month=10, day=1)
    house_reading = dictMeter_clean[key]
    reading_arr = []
    for i in range(24*183):
        sub_row = house_reading[(house_reading['localminute']>(start_time+pd.
        ↳Timedelta(hours=i))) & \
                                (house_reading['localminute']<(start_time+pd.
        ↳Timedelta(hours=i+1))))]
        hourly_reading = sub_row['meter_value'].mean()
        reading_arr.append(hourly_reading)
    hourly_readings[key] = reading_arr
```

[276]: *# transfer the array to dataframe*

```
df = pd.DataFrame(hourly_readings)
```

[277]: *# check the final completed data frame*

df

```
[277]:
```

	739	8890	6910	3635	1507	5810	\
0	88858.0	197164.0	179118.0	151322.0	390354.000	97506.0	
1	88858.0	197164.8	179118.0	151330.0	390354.000	97508.0	
2	88859.0	NaN	179118.0	151330.0	390354.000	97508.0	
3	88860.0	197166.0	179118.0	151330.0	390355.875	97508.0	
4	88860.0	197166.0	179118.0	151330.0	390356.000	97508.0	
...	
4387	103010.0	227306.0	NaN	170805.0	422680.000	115248.8	
4388	103011.0	NaN	205218.0	170806.0	422690.500	115254.0	
4389	103012.0	227308.0	205222.0	170806.0	422696.000	115270.0	
4390	103012.0	NaN	205224.0	170806.0	422704.000	115270.0	
4391	103012.0	227326.0	205224.0	170816.0	422708.000	115272.0	

	484	4352	1718	1714	...	8244	2755	9600	\
0	99298.000000	218216.0	161076.0	147048.000000	...	NaN	NaN	NaN	
1	99299.428571	218216.0	161076.0	147048.000000	...	NaN	NaN	NaN	
2	99300.000000	218218.0	161076.0	147048.000000	...	NaN	NaN	NaN	
3	99300.000000	NaN	161077.6	147051.142857	...	NaN	NaN	NaN	
4	99302.222222	218218.0	161078.0	147058.000000	...	NaN	NaN	NaN	
...	
4387	114168.500000	NaN	180958.0	170164.000000	...	NaN	NaN	NaN	
4388	114171.538462	NaN	180958.0	170164.000000	...	NaN	NaN	NaN	
4389	114172.000000	NaN	180966.0	170164.000000	...	NaN	NaN	NaN	
4390	114172.000000	NaN	180980.0	170164.000000	...	NaN	NaN	NaN	
4391	114172.666667	NaN	180980.0	170164.000000	...	NaN	NaN	NaN	

	2946	1403	7566	6673	2814	6101	4874
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...
4387	NaN	NaN	NaN	85354.0	NaN	NaN	NaN
4388	NaN	NaN	NaN	85356.0	NaN	NaN	NaN
4389	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4390	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4391	NaN	NaN	NaN	85362.0	NaN	NaN	NaN

[4392 rows x 157 columns]

```
[278]: # build the model and do the prediction using linear regression
```

```
time_arr = []  
start_time = pd.Timestamp(year=2015, month=10, day=1, hour=0)
```

```

for i in range(4392):
    time_arr.append(start_time + pd.Timedelta(hours=i))

hour_dict = {}

for key in df:
    reg = LinearRegression().fit(np.array(df[key].dropna().index).
    ↪reshape(-1,1), df[key].dropna())

    time_arr0 = time_arr
    index = df[key].index[df[key].apply(np.isnan)]
    index = list(index)
    time_arr0 = np.delete(time_arr, index)
    """
    fig, axs = plt.subplots(2, 1, figsize=(20,8))
    axs[0].plot(np.array(df[key].dropna().index).reshape(-1,1), \
                reg.predict(np.array(df[key].dropna().index).reshape(-1,1)), \
    ↪color='r')
    axs[0].scatter(range(4392), df[key], color='b', s=1)
    axs[0].set_title("Meter ID {} Average Hourly Meter Reading".format(key))

    axs[1].plot(time_arr0, \
                reg.predict(np.array(df[key].dropna().index).reshape(-1,1)), \
    ↪color='r')
    axs[1].scatter(dictMeter_clean[key]["localminute"], \
    ↪dictMeter_clean[key]["meter_value"], color='g', s=1)
    axs[1].set_title("Meter ID {} Actual Meter Reading".format(key))


    plt.tight_layout()
    axs[0].legend(["LR Fitting Line", "Average Hourly Reading"])
    axs[1].legend(["LR Fitting Line", "Actual Meter Reading"])
    plt.show()
    print("The accuracy score of fitting for meter ID {} is {}"\
          .format(key, reg.score(np.array(df[key].dropna().index).
    ↪reshape(-1,1), df[key].dropna()))))
    print("The next predicted average hourly reading for meter ID {} for the \
    ↪period 2016-04-01 00:00:00 to 01:00:00 is {}"\
          .format(key, reg.predict(np.array([[4392]]))[0]))
    print("The average hourly consumption for meter ID {} is {}".format(key, \
    ↪(reg.predict(np.array([[4393]])-reg.predict(np.array([[4392]])))[0]))
    \
    ↪print("-----")
    """
    hour_dict[key] = (reg.predict(np.array([[4393]]))-reg.predict(np.
    ↪array([[4392]])))[0]

```

```
[279]: # plot the predicted average hourly consumption for all homes

plt.figure(figsize=(100,20))
readingLst = [i for i in list(hour_dict.values())]
meterLst = [str(i) for i in list(hour_dict.keys())]
plt.bar(meterLst,readingLst,color="b",label="Predicted Average Hourly_
↳Consumption for each home using Linear Regression")
plt.xlabel("Meter ID")
plt.ylabel("Predicted Average Hourly Consumption")
plt.title("Predicted Average Hourly Consumption for each home")
```

```
[279]: Text(0.5, 1.0, 'Predicted Average Hourly Consumption for each home')
```



7 Discussion

We would like to divide all 157 homes into 4 categories based on the average monthly consumption. 30% of homes will fall in low consumption tier, 30% of homes will fall in medium consumption tier, 30% of homes will fall in high consumption tier and the rest 10% of homes will fall in extremely high consumption tier.

Low consumption tier 1:	0 - 3348
Medium consumption tier 2:	3348 - 4612.8
High consumption tier 3:	4612.8 - 6859.68
Extremely high consumption 4:	6859.68 and above

There are 47 homes fall in tier 1, 47 homes fall in tier 2, 47 homes fall in tier 3 and 16 homes fall in tier 4. Please refer to below for the details of homes meter ID for each tier.

To utilise the energy more efficiently and encourage families to save and reduce the waste of energy, The gas company can set plans of different price per cubic meter for different tiers from low to high.

For example, gas price for tier 1 is 0.5 dollar per cubic meter, 0.6 dollar per cubic meter for tier 2, 0.7 dollar per cubic meter for tier 3 and 1 dollar per cubic meter for tier 4. If your family consumption for this month is 4000, then you should pay $0.5 \times 3348 + 0.6 \times (4000 - 3348) = 2065.2$ dollars for this month.

7.1 To do in the future

1. As all the results are based on the linear model of prediction, we can improve our prediction to better fit the model. Please refer to below improvement part.
2. We should also consider the if one family has more members, the family is supposed to consume more gas per month, so the next step is to gather the information of family members for all the 157 families. Then we can get the more accurate data by calculating the average monthly consumption per person and set the price tiers.
3. Moreover, according to the monthly consumption level of each home, the construction company can build more proper and better pipeline when doing the maintenance or reconstruction.

```
[280]: # create the dataframe from the hourly consumption dictionary and sort the
      ↪ dataframe in ascending order

hour_df = pd.DataFrame.from_dict(hour_dict, orient='index', columns=['Hourly
      ↪ Reading'])
hour_df.sort_values(by='Hourly Reading', inplace=True)
hour_df
```

```
[280]:      Hourly Reading
8967      0.360368
4671      0.979432
5317      1.006564
9849      1.356889
77        1.369670
...
2755      11.428275
9639      12.714881
222       16.092826
1790      16.716645
2378      16.999217

[157 rows x 1 columns]
```

```
[281]: # separate the whole dataframe into 4 tiers according to the percentage

low_df = hour_df.iloc[0:47] # low 30% homes
med_df = hour_df.iloc[47:94] # medium 30% - 60% homes
high_df = hour_df.iloc[94:141] # high 60% - 90% homes
exhigh_df = hour_df.iloc[141:157] # extreme high 90% - 100% home
```

```
[282]: # get the 4 tiers consumption range

low_tier = [0, 4.5]
med_tier = [4.5, 6.2]
high_tier = [6.2, 9.22]
exhigh_tier = [9.22, 'above']
```

```
[283]: # count the number of families fall in each of tiers


low_count = low_df.count()['Hourly Reading']
med_count = med_df.count()['Hourly Reading']
high_count = high_df.count()['Hourly Reading']
exhigh_count = exhigh_df.count()['Hourly Reading']

[284]: # get each tier consumption range for a month assuming 31 days in a month

low_tier31 = [low_tier[0]*24*31, low_tier[1]*24*31]
med_tier31 = [med_tier[0]*24*31, med_tier[1]*24*31]
high_tier31 = [high_tier[0]*24*31, high_tier[1]*24*31]
exhigh_tier31 = [exhigh_tier[0]*24*31, exhigh_tier[1]]

[285]: # plot the bar chart for the 4 tiers and get all the homes meter ID for each of
↳ tiers

x_axis = ['Low 0-3348', 'Medium 3348-4612.8', 'High 4612.8-6859.68', 'Extreme_
↳ high 6859.68 and above']
y_axis = [low_count, med_count, high_count, exhigh_count]
plt.figure(figsize=(10,5))
plt.bar(x_axis, y_axis)
plt.xlabel('Average monthly consumption ')
plt.ylabel('Number of homes fall in the tier')
plt.tight_layout()
plt.show()
print('Number of homes fall in low consumption tier is {}, the meter ID is {}'.
↳ format(low_count, list(low_df.index)))
print('-----')
print('Number of homes fall in medium consumption tier is {}, the meter ID is_
↳ {}'.format(med_count, list(med_df.index)))
print('-----')
print('Number of homes fall in high consumption tier is {}, the meter ID is {}'.
↳ format(high_count, list(high_df.index)))
print('-----')
print('Number of homes fall in extreme high consumption tier is {}, the meter_
↳ ID is {}'.format(exhigh_count, list(exhigh_df.index)))
```



Number of homes fall in low consumption tier is 47, the meter ID is [8967, 4671, 5317, 9849, 77, 1792, 5658, 1403, 8467, 6578, 3544, 6673, 1791, 44, 9956, 2645, 35, 8084, 4421, 2980, 4228, 6836, 1619, 5403, 4447, 7739, 5129, 370, 2034, 5814, 8059, 9982, 8386, 484, 3778, 739, 2449, 3039, 2965, 1042, 8703, 9278, 3849, 1800, 3134, 7989, 9729]

Number of homes fall in medium consumption tier is 47, the meter ID is [3893, 9631, 5439, 871, 5810, 114, 6830, 4296, 7287, 7429, 4031, 1556, 4732, 7965, 3310, 7030, 5131, 8156, 1697, 1801, 7741, 7117, 6685, 2461, 5275, 4352, 7682, 2945, 1718, 2470, 1415, 5395, 3635, 4029, 5193, 5636, 8829, 8155, 1185, 5892, 2094, 1086, 5484, 9766, 5545, 1283, 252]

Number of homes fall in high consumption tier is 47, the meter ID is [1103, 2072, 744, 7919, 8086, 187, 2575, 2638, 94, 6412, 3527, 483, 1714, 6505, 4356, 3367, 2818, 9121, 5785, 4767, 4373, 3723, 6101, 9295, 7900, 661, 6910, 2814, 6863, 7674, 9474, 2233, 4193, 4998, 4514, 3577, 9600, 2018, 8244, 7566, 2129, 8890, 9052, 4874, 2335, 1507, 7016]

Number of homes fall in extreme high consumption tier is 16, the meter ID is [3036, 7460, 3918, 7017, 9160, 1589, 9620, 7794, 5972, 9134, 2946, 2755, 9639, 222, 1790, 2378]

8 Model Improvement Part

Below is our improved model to predict the daily consumption across the 6 months for all the individual homes. We used **LSTM (long short-term memory)** model. We showed the first 10

homes meter for demonstration.

```
[296]: import pandas as pd
from datetime import datetime, date
import numpy as np
import matplotlib.pyplot as plt
import torch.nn as nn
import torch
from sklearn.preprocessing import MinMaxScaler

# LSTM model used in our experiments with a single LSTM layer and a linear
↳ layer.
class LSTM(nn.Module):
    def __init__(self, input_size, hidden_layer_size, output_size):
        super().__init__()
        self.hidden_layer_size = hidden_layer_size

        self.lstm = nn.LSTM(input_size, hidden_layer_size)

        self.linear = nn.Linear(hidden_layer_size, output_size)

        self.hidden_cell = (torch.zeros(1, 1, self.hidden_layer_size),
                             torch.zeros(1, 1, self.hidden_layer_size))

    def forward(self, input_seq):
        lstm_out, self.hidden_cell = self.lstm(input_seq.view(len(input_seq),
↳ 1, -1), self.hidden_cell)
        predictions = self.linear(lstm_out.view(len(input_seq), -1))
        return predictions[-1]

# use the raw daily consumption data to build input data, with the first
↳ element contain list of tw items corresponding to data in tw days, the
↳ second tuple element contains data in the tw+1st day
def create_input_sequences(input_data, tw):
    inout_seq = []
    L = len(input_data)
    for i in range(L-tw):
        train_seq = input_data[i:i + tw]
        train_label = input_data[i + tw:i + tw + 1]
        inout_seq.append((train_seq, train_label))
    return inout_seq

def train(epochs, model, train_inout_seq):
    # define loss function and optimizer
```

```

loss_function = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
pred = []
for i in range(epochs):
    model.train()
    for seq, labels in train_inout_seq:
        seq = seq
        labels = labels
        optimizer.zero_grad()
        # reset model hidden states
        model.hidden_cell = (torch.zeros(1, 1, model.hidden_layer_size),
                             torch.zeros(1, 1, model.hidden_layer_size))

        y_pred = model(seq)
        if i == epochs - 1:
            pred.append(y_pred.item())
            single_loss = loss_function(y_pred, labels)
            single_loss.backward()
            optimizer.step()

    if i % 25 == 1:
        print(f'epoch: {i:3} loss: {single_loss.item():10.8f}')
print(f'epoch: {i:3} loss: {single_loss.item():10.10f}')
return pred

def test(model, test_data, test_len, tw):
    # predict future consumption according to known consumption
    model.eval()
    model = model
    while len(test_data) < test_len:
        seq = torch.FloatTensor(test_data[-tw:])
        with torch.no_grad():
            # reset model hidden states
            model.hidden = (torch.zeros(1, 1, model.hidden_layer_size),
                             torch.zeros(1, 1, model.hidden_layer_size))
            test_data.append(model(seq).item())
    return test_data

def read_data(data_path, unique_house):
    # read daily consumption data corresponding to a single house
    daily_reading = np.zeros(183)
    csv_data = pd.read_csv(data_path,
                           date_parser=lambda x: datetime.strptime(x[:19],
                               ↪ '%Y-%m-%d %H:%M:%S'),
                           parse_dates=['localminute'])

```

```

start_time = pd.Timestamp(year=2015, month=10, day=1)
house_reading = csv_data[csv_data['dataid'] == unique_house].\
    sort_values(by='localminute', ascending=True).reset_index()
for i in range(183):
    start_row = house_reading[house_reading['localminute'] < start_time +
→pd.Timedelta(days=i)].reset_index()
    if len(start_row) > 0:
        start_row = start_row.iloc[-1]
    else:
        continue
    end_row = house_reading[house_reading['localminute'] < start_time + pd.
→Timedelta(days=i+1)].reset_index()
    if len(end_row) > 0:
        end_row = end_row.iloc[-1]
    else:
        continue
    daily_reading[i] += end_row['meter_value'] - start_row['meter_value']

return daily_reading

def main():
    data_path = 'dataport-export_gas_oct2015-mar2016.csv'
    csv_data = pd.read_csv(data_path,
                           date_parser=lambda x: datetime.strptime(x[:19],
→'%Y-%m-%d %H:%M:%S'),
                           parse_dates=['localminute'])
    unique_houses = csv_data['dataid'].unique()
    for unique_house in unique_houses[:10]:
        train_data = read_data(data_path, unique_house)

        # normalize data with minmaxscaler
        scaler = MinMaxScaler(feature_range=(-1, 1))
        train_data_transformed = scaler.fit_transform(train_data.reshape(-1, 1))
        train_data_transformed = torch.FloatTensor(train_data_transformed).
→view(-1)

        # input sequence length is set to 30, since we are using daily
→consumption data, and there are about 30 days in a month
        train_window = 30
        train_seq = create_input_sequences(train_data_transformed, train_window)

        model = LSTM(input_size=1,
                      hidden_layer_size=256,
                      output_size=1)
        pred = train(150, model, train_seq)
        real = []

```

```

    for item in train_seq:
        real.append(item[1].item())

    # recover real data from normalized data
    real = scaler.inverse_transform(np.array(real).reshape(-1, 1))[:, 0].
    ↪tolist()
    pred = scaler.inverse_transform(np.array(pred).reshape(-1, 1))[:, 0].
    ↪tolist()

    # plot real and predicted daily consumption
    x = [i for i in range(len(pred))]
    plt.plot(x, pred, color='r', label='pred')
    plt.plot(x, real, color='b', label='real')
    plt.xlabel('day')
    plt.ylabel('daily consumption')
    plt.title(str(unique_house))
    plt.legend()
    plt.show()


if __name__ == '__main__':
    main()

```


```

epoch:   1 loss: 0.04398124
epoch:  26 loss: 0.00025985
epoch:  51 loss: 0.00650813
epoch:  76 loss: 0.00255294
epoch: 101 loss: 0.00184251
epoch: 126 loss: 0.00090139
epoch: 149 loss: 0.0000961946


```




```
epoch: 1 loss: 0.02326862
epoch: 26 loss: 0.00620577
epoch: 51 loss: 0.00000001
epoch: 76 loss: 0.01310374
epoch: 101 loss: 0.01540042
epoch: 126 loss: 0.00045660
epoch: 149 loss: 0.0052261115
```


```
epoch: 1 loss: 0.04420021
epoch: 26 loss: 0.00038037
epoch: 51 loss: 0.00000284
epoch: 76 loss: 0.00329855
epoch: 101 loss: 0.02685077
epoch: 126 loss: 0.00368978
epoch: 149 loss: 0.0044701588
```




```
epoch: 1 loss: 0.00156483
epoch: 26 loss: 0.00000141
epoch: 51 loss: 0.00416336
epoch: 76 loss: 0.00542395
epoch: 101 loss: 0.00330202
epoch: 126 loss: 0.00269620
epoch: 149 loss: 0.0024086898
```




epoch: 1 loss: 0.05147930
epoch: 26 loss: 0.00655148
epoch: 51 loss: 0.01267585
epoch: 76 loss: 0.00021046
epoch: 101 loss: 0.00519592
epoch: 126 loss: 0.00925888
epoch: 149 loss: 0.0038447215




```
epoch: 1 loss: 0.00984569
epoch: 26 loss: 0.02209531
epoch: 51 loss: 0.06282237
epoch: 76 loss: 0.04155124
epoch: 101 loss: 0.04403833
epoch: 126 loss: 0.02124080
epoch: 149 loss: 0.0222013518
```




epoch: 1 loss: 0.10708092
epoch: 26 loss: 0.06553694
epoch: 51 loss: 0.06468331
epoch: 76 loss: 0.05946452
epoch: 101 loss: 0.07530762
epoch: 126 loss: 0.10952046
epoch: 149 loss: 0.0145193953




```
epoch: 1 loss: 0.00002735
epoch: 26 loss: 0.00046538
epoch: 51 loss: 0.00000478
epoch: 76 loss: 0.00000025
epoch: 101 loss: 0.00004230
epoch: 126 loss: 0.00012192
epoch: 149 loss: 0.0002039582
```



```
epoch: 1 loss: 0.00214029
epoch: 26 loss: 0.00140670
epoch: 51 loss: 0.00041599
epoch: 76 loss: 0.00090980
epoch: 101 loss: 0.00213046
epoch: 126 loss: 0.00167145
epoch: 149 loss: 0.0030836815
```



epoch: 1 loss: 0.06962930
epoch: 26 loss: 0.00007665
epoch: 51 loss: 0.05566140
epoch: 76 loss: 0.07088415
epoch: 101 loss: 0.00790488
epoch: 126 loss: 0.00674224
epoch: 149 loss: 0.0000206286



[]: